

THEJUS MAHAJAN PH.D.

Computational Scientist | Machine Learning · Interpretability · Clinical Data

+49 1525 9697629

@ thejus.mahajan@uni-hamburg.de

thejusmahajan.github.io

Reventlowstraße 17, 22605 Hamburg

*30.09.1992 (Talikulam, India)

Married

Nationality: Indian

Arbeitsberechtigung für Deutschland vorhanden · Work authorisation: yes

linkedin.com/in/thejusmahajan

github.com/thejusmahajan

Google Scholar



PROFILE

- Computational Scientist (PhD, Post-doc) working at the intersection of machine learning and scientific computing: 10+ years building and validating computational models across physics, ocean biogeochemistry, bioinformatics and clinical data.
- Independent research in **neural network interpretability**: extracting and correcting internal attention representations from a 15-layer transformer (Leela Chess Zero BT3) in PyTorch/ONNX. Write-up: thejusmahajan.github.io/blog-1c0-attention-frame.html
- Core skills: Python (PyTorch, ONNX, JAX, NumPy, Pandas), R (tidyverse, Shiny, tidymodels), SQL, Git, HPC.
- HealthTwiSt GmbH (Praxisphase): refactored a production R pipeline for the DeGIR quality registry (143K+ patient records, ~300 clinics), built an R Shiny dashboard, created the degirtools R package.
- Available immediately. Open to: Machine Learning / AI Engineering, Research Software Engineering, Data Science, Computational Biology; open to relocation within Germany.

PROFESSIONAL EXPERIENCE

Independent Research — Neural Network Interpretability

Self-directed (open source)

2026 - present

Hamburg, Germany

Mechanistic interpretability of a 15-layer transformer (Leela Chess Zero BT3, 24 heads/layer)

- Built an activation-capture toolchain in **PyTorch**: converted the network from **ONNX** to a PyTorch module tree and registered forward hooks on all 15 encoder query-key softmax modules to capture per-head 64×64 attention tensors during inference.
- Diagnosed a **reference-frame defect** that silently mirrored every attention map for half of all positions — no exception, no failing test, plausible output. Identified quantitatively (mean per-square residual 0.0003 against the rank-flipped map over 40 positions) and fixed behind an explicit API boundary with regression tests.
- Found a second, larger defect while writing the first up: the model's 112-plane input includes ~84 history planes that were being left empty, so every forward pass ran out of distribution. Validated the fix against the reference engine as ground truth — top-1 policy agreement rose from 1/6 to 5/6.
- Engineered **batched GPU/CPU inference** and an async engine-orchestration layer (UCI process management, raw policy-head extraction before search, position-level parallelism).
- Authored a normative metrics module (710 lines, pure functions) and a deterministic symbolic feature extractor (787 lines) to ground model outputs in verifiable structure rather than generated text.
- Delivered end to end: ~20,400 lines of Python (FastAPI) and ~6,700 of TypeScript (React 19), covered by 339 automated tests (290 backend, 49 frontend).

Clinical Data Engineering — Praxisphase (Weiterbildung ABI-2025-3)

HealthTwiSt GmbH

02/2026 - 04/2026

Berlin, Germany

Clinical data research and development under Dr. Andreas Busjahn

- Refactored a monolithic 1,834-line R data pipeline (tidyverse) for Germany's interventional radiology quality registry (DeGIR); 143,000+ patient records from ~300 clinics. Reduced to 1,349 lines (−26.5%) with byte-identical output verified via `identical()`.

KEY STRENGTHS



Neural Network Interpretability & Deep Learning

Attention extraction from a 15-layer transformer via PyTorch forward hooks on ONNX-converted weights. Diagnosed and fixed two silent correctness defects in the model-analysis pipeline; batched GPU inference.



Clinical Data Pipeline Engineering

Refactored a production data pipeline (1,834 lines, 143K records, ~300 clinics) with 26.5% code reduction. Built degirtools R package. Byte-level verification after every change via `identical()`.



Biostatistics & Machine Learning

Proficient in multivariate analysis (PCA, ANOVA), clinical outcome analysis, developing machine learning workflows using tidymodels, and cross-validation.



Bioinformatics & Reproducible Research

Proficient in NGS data analysis, building reproducible pipelines with Nextflow (DSL2), Bioconductor workflows, and ensuring GDPR-compliant data handling.

TECHNICAL SKILLS

Programming

Python (PyTorch, ONNX, JAX, NumPy, Pandas, FastAPI), R (tidyverse, Shiny, tidymodels, Bioconductor), SQL, TypeScript, Fortran, C++, BASH

Machine Learning

Transformers, attention/activation capture (forward hooks), mechanistic interpretability, ONNX→PyTorch conversion, batched GPU inference, policy/value head analysis, tidymodels

Clinical Data & Biostatistics

Medical registry pipelines, GDPR/DSGVO compliance, multivariate analysis, PCA, ANOVA, survival analysis concepts, clinical endpoints

Data Engineering

Pipeline refactoring, config-driven ETL, `identical()` verification, R Shiny dashboards, large datasets (HDF5, NetCDF)

Bioinformatics

NGS analysis, Nextflow (DSL2), sequence analysis, Bioconductor

Tools

Git/GitHub, Linux/Unix, HPC clusters, Conda, LaTeX, Docker concepts

- Externalised 257 hard-coded correction rules into 8 configuration CSV files, enabling non-programmer editing of business rules.
- Replaced 360 deprecated function calls across 4 pipeline files, modernising the codebase to current tidyverse standards.
- Created the `degirtools` R package (9 exported functions, roxygen2 docs, `devtools::check()` passing) by extracting reusable logic from a 767-line helper file.
- Discovered 2 pre-existing bugs through systematic code analysis; documented both and preserved original behaviour for supervisor review — demonstrating production-code discipline.
- Built an interactive R Shiny dashboard (6 modules: overview, interventions, complications, radiation doses, success rates, about) using plotly, DT, and bslib. Deployed to shinyapps.io with GDPR-safe synthetic data.
- Wrote Tidymodels teaching materials (Quarto/GitHub) replacing the legacy caret framework for ML workflows in R.

Continuing Education – Bioinformatics and Biostatistics

CQ Beratung + Bildung GmbH

📅 08/2025 - 02/2026 📍 Berlin, Germany

Applied bioinformatics and biostatistics

- Programming methodology: Shell (BASH), Python (object-oriented), R.
- Bioinformatics databases (NCBI, Biopython, SQL) and sequence analysis.
- NGS analysis pipelines (Next-Generation Sequencing; Nextflow, Galaxy); structural bioinformatics.
- Biostatistics (R/Bioconductor): ANOVA, PCA, HCA, multivariate analysis, Markov chains.
- Capstone: Built a reproducible Nextflow (DSL2) pipeline for Hepatitis Delta Virus sequence analysis — automated NCBI download, MAFFT alignment, and trimAl trimming with Conda-managed dependencies (🔗 GitHub).

Guest Scientist (parallel to continuing education from 08/2025)

Helmholtz-Zentrum Hereon, Ecosystem Modelling

📅 05/2025 - 10/2025 📍 Geesthacht, Germany

Modelling evolutionary processes in marine ecosystems; simulation of long-term environmental impacts using advanced computational methods. Finalised the computational modelling framework and prepared the simulation data for peer-reviewed publication.

Job Search & German Language Course

📅 02/2025 - 04/2025 📍 Hamburg, Germany

Active job search and German language acquisition (Goethe-Zertifikat B1).

Post-doctoral Scientist

Universität Hamburg, Institute for Marine Ecosystem and Fishery Science

📅 08/2021 - 01/2025 📍 Hamburg, Germany

Marine ecosystem modelling

- *Parental leave from 01/2024 to 12/2024.*
- Model development: Development of the “Cyanobacteria Life Cycle” (CLC) model within the ERGOM framework for climate warming projections.
- High-performance computing (HPC): Translation of a legacy Fortran model into a high-performance Python engine (**Google JAX**, TPU/GPU parallelisation). Processing of large NetCDF datasets in Linux HPC environment.

Career Transition

📅 04/2021 - 07/2021 📍 Hamburg, Germany

Relocation from India to Germany and active job search.

Physics Tutor & Chess Coach — Freelance / Tutorwaves Solutions Inc., Kerala, India · 10/2018 – 03/2021

Doctoral Researcher in Astrochemistry

LANGUAGES

English



Fluent (C1)

German



B1 (Goethe-Zertifikat), B2 exam scheduled

French



Basic knowledge

Malayalam



Native

Tamil



Fluent

Hindi



Basic knowledge

COURSES & TRAINING

High Performance Computing (HPC) Training

Jülich Supercomputing Center (JSC),
Forschungszentrum Jülich

MITx: 7.00x Introduction to Biology – The Secret of Life

Comprehensive coursework in genetics, biochemistry, and molecular biology.

PUBLICATIONS (ASTROCHEMISTRY / PHYSICS, 2018–2020)

T. Mahajan et al. *J. Phys.: Conf. Ser.*, 1412:142026, 2020.

T. Idbarkach et al. *Astron. Astrophys.*, 628:A75, 2019.

T. Mahajan et al. *J. Phys. B*, 2019.

T. Idbarkach et al. *J. Phys. B*, 51(24):245201, 2018.

T. Idbarkach et al. *J. Phys.: Conf. Ser.*, 1412(11):112028, 2020.

🔗 Google Scholar Profile

Conducting atomic collision experiments, data analysis, and results modelling.

- Acquisition and processing of terabytes of raw data from tandem particle accelerator experiments.
- C++ code optimisation: improvement of existing algorithms for noise suppression, signal smoothing, and conversion of raw data into multi-dimensional structures.
- Development of a molecular fragmentation model for predicting chemical reactions in space (contribution to the KIDA database).

EDUCATION

Ph.D. in Astrochemistry

Université Paris-Saclay, Institute of Molecular Sciences (ISMO)

10/2015 - 09/2018 Orsay, France

Supervisor: Dr. Karine Béroff. Thesis: *Excitation and fragmentation of C_nN^+ ($n=1-3$) in collisions with He atoms at intermediate velocity.*

M.Sc. in Physics

National Institute of Technology Calicut, India

07/2012 - 12/2014 Calicut, India

First Class with Distinction (CGPA 8.71/10).

Research Project – Theoretical Atomic Physics

IIT Mandi (under Dr. Hari Varma)

02/2015 - 08/2015 Mandi, India

Theoretical atomic physics; partly remote.

B.Sc. in Physics

University of Calicut, India

06/2009 - 04/2012 Thrissur, India

Grade: B+ (CGPA 3.49/4.0).

Hamburg, 18 April 2026



Dr. Thejus Mahajan